

The Method of Moderation

Christopher D. Carroll^{a,b}, Alan Lujan^{a,b,*}, Karsten Chipeniuk^c, Kiichi Tokuoka^d, Weifeng Wu^e

^a*Johns Hopkins University, Krieger School of Arts and Sciences Baltimore, MD,*

^b*Econ-ARK,*

^c*Reserve Bank of New Zealand,*

^d*Japanese Ministry of Finance,*

^e*Fannie Mae, Washington, DC,*

Abstract

In a risky world, a pessimist assumes the worst will happen. Someone who ignores risk altogether is an optimist. Consumption decisions are mathematically simple for both the pessimist and the optimist because both behave as if they live in a riskless world. A consumer who is a realist (that is, who wants to respond optimally to risk) faces a much more difficult problem, but (under standard conditions) will choose a level of spending somewhere between that of the pessimist and the optimist. We use this fact to redefine the space in which the realist searches for optimal consumption rules. The resulting solution accurately represents the numerical consumption rule over the entire interval of feasible wealth values with remarkably few computations.

Keywords: Dynamic Stochastic Optimization, Consumption-Saving Models, Numerical Methods

JEL classification: D14; C61; G11

1. Introduction

Solving a consumption-saving problem using numerical methods requires the modeler to choose how to represent a policy function. In the stochastic case, where analytical solutions are generally not available, a common approach is to use low-order polynomial splines that exactly match the function at a finite set of gridpoints, and then to assume that interpolated or extrapolated versions of that spline represent the function well at the continuous infinity of unmatched points. [Carroll \(2006\)](#) developed the endogenous gridpoints method (EGM), which has become a standard tool for computing consumption at gridpoints determined endogenously using the Euler equation.

*Corresponding author

Email addresses: ccarroll1@jhu.edu (Christopher D. Carroll), alujan@jhu.edu (Alan Lujan)

Unfortunately, this endogenous gridpoints solution is not very well-behaved outside the original range of gridpoints, though other common solution methods are no better outside their own predefined ranges. Figure 1 demonstrates the point. The figure shows the approximated precautionary component of saving, the amount by which the realist consumes less than an optimist with the same expected income path. Theory proves that precautionary saving is always positive, yet the linearly extrapolated numerical approximation eventually predicts negative precautionary saving. Under uncertainty, however, the consumption-saving rule must be evaluated outside *any* prespecified grid, because large positive shocks push a sufficiently wealthy individual's next-period assets beyond the grid boundaries.




Figure 1: For Large Enough Resources m , Predicted Precautionary Saving is Negative (Oops!)

This error cannot be fixed by extending the upper gridpoint. While extrapolation techniques can prevent this from being fatal, the problem is often dealt with using inelegant methods whose implications for accuracy are difficult to gauge.

This paper argues that, in the standard consumption problem, a better approach is to rely upon the fact that without uncertainty, the optimal consumption function has a simple analytical solution. The key insight is that, under standard assumptions, the consumer who faces an uninsurable labor income risk will consume less than a consumer with the same path for expected income but who does not perceive any uncertainty. Following [Leland \(1968\)](#), [Sandmo \(1970\)](#), and [Kimball \(1990\)](#), the ‘realist’ consumer, who *does* perceive the risks, will engage in ‘precautionary saving,’ so the perfect foresight riskless solution provides an upper bound to the solution that will actually be optimal. A lower bound is provided by the behavior of a consumer who has the subjective belief that the future level of income will be the worst that it can possibly be. This consumer, too, behaves according to the convenient analytical perfect foresight solution, but his certainty is that of a pessimist perfectly confident in his pessimism.

We build on bounds for the consumption function and limiting MPCs established by [Stachurski and Toda \(2019\)](#), [Ma et al. \(2020\)](#), [Carroll \(2009\)](#), and [Ma and Toda \(2021\)](#) in buffer-stock theory. Using results from [Carroll and Shanker \(2024\)](#), we show how to use these bounds to constrain the shape and characteristics of the solution to the ‘realist’ problem characterized by [Carroll \(1997\)](#). Imposition of these constraints clarifies and speeds the solution of the realist’s problem. For comparison, we use the endogenous gridpoints method ([Carroll, 2006](#)) as our benchmark, which computes consumption at gridpoints determined endogenously using the Euler equation.

After showing how to use the method in the baseline case, we show how to refine it to encompass an even tighter theoretical bound, and how to extend it to solve a problem in which the consumer faces both labor income risk and rate-of-return risk.

2. The Realist’s Problem

Consider a consumer who correctly perceives all risks. The consumer has CRRA utility over consumption c with risk aversion parameter $\rho > 0$:

$$\mathbf{u}(c) = \begin{cases} \frac{c^{1-\rho}}{1-\rho} & \text{if } \rho \neq 1 \\ \log c & \text{if } \rho = 1. \end{cases} \quad (1)$$

This utility function satisfies prudence ($\mathbf{u}''' > 0$), which induces precautionary saving. The consumer maximizes expected lifetime utility, discounted by the time preference factor β :

$$\max_{c_t, a_t} \mathbb{E}_t \left[\sum_{n=0}^{T-t} \beta^n \mathbf{u}(c_{t+n}) \right] \quad (2)$$

subject to $c_t + a_t = m_t$, where m denotes market resources and a denotes assets. We focus on resources of the form $m_{t+1} = a_t R_{t+1} + y_{t+1}$, where R_{t+1} denotes the interest rate, and y_{t+1} labor income. Initially we take R_{t+1} to be deterministic, and relax this later.

While our method can be adapted to a range of stochastic labor income processes, to fix ideas we suppose income evolves via the Friedman-Muth process ([Friedman \(1957\)](#) distinguished permanent from transitory income; [Muth \(1960\)](#) provided the stochastic framework). That is, $y_{t+1} = p_{t+1} \xi_{t+1}$ where p denotes permanent labor income and ξ_{t+1} a transitory component. Permanent income growth is given by $p_{t+1} = p_t \mathcal{G}_{t+1}$, where the gross growth factor $\mathcal{G}_{t+1} = G_{t+1} \psi_{t+1}$. Here G_{t+1} is the deterministic permanent income growth factor, while ψ_{t+1} are permanent shocks with mean unity and bounded support $[\underline{\psi}, \bar{\psi}]$ where $0 < \underline{\psi} \leq 1 \leq \bar{\psi} < \infty$. Transitory shocks ξ_{t+1} take value 0 with probability $\varphi > 0$ (unemployment) or $\theta_{t+1}/(1-\varphi)$ otherwise, with $\mathbb{E}_t[\theta_{t+1}] = 1$.

This problem can be rewritten (see [Carroll \(2020b\)](#) for a proof) in a more convenient form in which choice and state variables are normalized by the level

of permanent income, e.g., replacing m_t with m_t/p_t . When that is done, the Bellman equation for the value function $\mathbf{v}$ of the transformed problem is

$$\mathbf{v}_t(m_t) = \max_{c_t, a_t} \left(\mathbf{u}(c_t) + \beta \mathbb{E}_t[\mathcal{G}_{t+1}^{1-\rho} \mathbf{v}_{t+1}(m_{t+1})] \right) \quad (3)$$

subject to the normalized transition $m_{t+1} = (R/\mathcal{G}_{t+1})a_t + \xi_{t+1}$, with Euler equation $\mathbf{u}'(c_t) = \beta R \mathbb{E}_t[\mathcal{G}_{t+1}^{-\rho} \mathbf{u}'(c_{t+1})]$.

Carroll and Shanker (2024) gives conditions for a finite solution of the problem with a Friedman-Muth process. Consider the case of time-invariant G_t , ψ_t , and R_t , and define the absolute patience factor $\Phi \equiv (\beta R)^{1/\rho}$. Then a finite solution requires: (i) finite-value-of-autarky condition (FVAC) $0 < \beta G^{1-\rho} \mathbb{E}[\psi^{1-\rho}] < 1$, (ii) absolute-impatience condition (AIC) $\Phi < 1$, (iii) return-impatience condition (RIC) $\Phi/R < 1$, (iv) growth-impatience condition (GIC) $\Phi/G < 1$, and (v) finite-human-wealth condition (FHWC) $G/R < 1$. These patience conditions ensure the consumption bounds and limiting MPCs used in our method.

For expositional simplicity, in the numerical results that follow, we assume $\rho \neq 1$ and set $G = 1$, $\psi = 1$ (no permanent income growth or shocks). The figures display the next-to-last period of a finite horizon problem, where the last-period analytical solution provides exact boundary conditions. The method extends naturally to the infinite-horizon case and to general parameter configurations.

3. The Method of Moderation

3.1. The Optimist, the Pessimist, and the Realist

As a preliminary to our solution, define $\bar{h}^1$ as end-of-period human wealth (the present discounted value of future labor income) for a perfect foresight version of the problem of a ‘risk optimist:’ a consumer who believes with perfect confidence that the shocks will always take their expected value of 1, $\xi_{t+n} = \mathbb{E}[\xi] = 1 \forall n > 0$. The solution to a perfect foresight problem of this kind takes the form

$$\bar{\mathbf{c}}(m) = (m + \bar{h})\underline{\kappa} \quad (4)$$

for a constant minimal marginal propensity to consume $\underline{\kappa}$.² We similarly define $\underline{h}$ ³ as ‘minimal human wealth,’ the present discounted value of labor income if the shocks were to take on their worst value in every future period $\xi_{t+n} = \underline{\xi} \forall n > 0$. We refer to a consumer who expects to encounter this sequence of shocks as a ‘pessimist’. Their consumption decision rule is given by

¹Setting $\xi_{t+n} = 1$ (the optimist’s assumption), human wealth in infinite-horizon is $\bar{h} = G/(R - G)$ if $R > G$. When $G = 1$, $\bar{h} = 1/(R - 1)$.

²The MPC of the perfect foresight consumer: infinite-horizon $\underline{\kappa} = 1 - \Phi/R$.

³Setting $\xi_{t+n} = \underline{\xi} \forall n > 0$, minimal human wealth is $\underline{h} = \underline{\xi}G/(R - G)$ if $R > G$. When $\underline{\xi} = 0$, $\underline{h} = 0$.

$$\underline{c}(m) = (m + \underline{h})\underline{\kappa}. \quad (5)$$

We will call a ‘realist’ the consumer who correctly perceives the true probabilities of the future risks and optimizes accordingly.

A first useful point is that, for the realist, a lower bound for the level of market resources is the natural borrowing constraint $\underline{m} = -\underline{h}$ derived by [Aiya-gari \(1994\)](#) and [Huggett \(1993\)](#), because if m equalled this value then there would be a positive finite chance (however small) of receiving $\xi_{t+n} = \underline{\xi}$ in every future period, which would require the consumer to set c to zero in order to guarantee that the intertemporal budget constraint holds. Since consumption of zero yields infinite marginal utility, [Zeldes \(1989\)](#) and [Deaton \(1991\)](#) show that the solution to the realist consumer’s problem is not well defined for values of $m \leq \underline{m}$, and the limiting value of the realist’s c is zero as $m \downarrow \underline{m}$ (established in [Carroll and Shanker \(2024\)](#)).

It is convenient to define ‘excess’ market resources as the amount by which actual resources exceed the lower bound, and ‘excess’ human wealth as the amount by which mean expected human wealth exceeds guaranteed minimum human wealth:

$$\begin{aligned} \Delta m &= m + \overbrace{\underline{h}}^{=-\underline{m}} \\ \Delta h &= \bar{h} - \underline{h}. \end{aligned} \quad (6)$$

We now rewrite the optimal consumption rules for the two perfect foresight problems in terms of excess resources and human wealth. The ‘pessimist’ perceives human wealth to be equal to its minimum feasible value $\underline{h}$ with certainty, and so consumes a constant fraction of current excess resources

$$\underline{c}(m) = \Delta m \underline{\kappa}. \quad (7)$$

The ‘optimist,’ on the other hand, pretends that there is no uncertainty about future income, and therefore consumes the same fraction out of current excess resources *and* excess human wealth

$$\bar{c}(m) = (\Delta m + \Delta h) \underline{\kappa}. \quad (8)$$

3.2. The Moderation Ratio

The pessimist expects the worst possible income in every future period with certainty, and must finance all future consumption through saving; this generates the most aggressive saving behavior. A realist, by contrast, can reoptimize as uncertainty resolves each period, so need not prepare for the worst with certainty. At the same time, the adverse outcome remains possible, so even a wealthy realist maintains some precautionary saving: a sufficiently well-off individual is nearly (but never completely) self-insured, and thus mostly smooths consumption like the optimist. The realist therefore consumes strictly more




Figure 2: Moderation Illustrated: $\underline{c} < \hat{c} < \bar{c}$

than the pessimist but strictly less than the optimist at every wealth level, as shown in Figure 2.

The proof is more difficult than might be imagined, but the necessary work is done in [Carroll and Shanker \(2024\)](#) so we will take the proposition as a fact:

$$\underline{c}(\underline{m} + \Delta m) < \hat{c}(\underline{m} + \Delta m) < \bar{c}(\underline{m} + \Delta m) \quad (9)$$

Subtracting $\underline{c}(\underline{m} + \Delta m)$ in each of these inequalities and using equations (7) and (8) gives

$$0 < \frac{\hat{c}(\underline{m} + \Delta m) - \underline{c}(\underline{m} + \Delta m)}{\underbrace{\left(\frac{\hat{c}(\underline{m} + \Delta m) - \underline{c}(\underline{m} + \Delta m)}{\Delta h \underline{\kappa}} \right)}_{\equiv \omega}} < \Delta h \underline{\kappa} \quad (10)$$

where the fraction in the middle of the last inequality is the moderation ratio measuring how close the realist's consumption is to the optimist's behavior (the numerator is the gap between the realist and pessimist) relative to the maximum possible gap between optimist and pessimist. When $\omega = 0$, the realist behaves like the pessimist (maximum precautionary saving); when $\omega = 1$, the realist behaves like the optimist (no precautionary saving). [Carroll and Kimball \(1996\)](#) and [Carroll and Shanker \(2024\)](#) establish that under bounded shocks, $\omega \in (0, 1)$ strictly for all $m > \underline{m}$. Defining $\mu = \log \Delta m$ (which can range from $-\infty$ to ∞), the object in the middle of the last inequality is

$$\omega(\mu) \equiv \left(\frac{\hat{c}(\underline{m} + e^\mu) - \underline{c}(\underline{m} + e^\mu)}{\Delta h \underline{\kappa}} \right), \quad (11)$$

and we now define

$$\begin{aligned}\chi(\mu) &= \log \left(\frac{\omega(\mu)}{1 - \omega(\mu)} \right) \\ &= \log(\omega(\mu)) - \log(1 - \omega(\mu))\end{aligned}\tag{12}$$

which has the virtue that it is *asymptotically linear* in the limit as μ approaches $+\infty$.⁴ Since $\omega \in (0, 1)$, the ratio $\omega/(1 - \omega)$ is the odds ratio, and χ is the log odds ratio, the same transformation that underpins logit regression in econometrics. The logit maps $\omega \in (0, 1)$ to $\chi \in (-\infty, \infty)$ with inverse sigmoid $\omega = 1/(1 + \exp(-\chi))$; log maps $(m - \underline{m}) \in (0, \infty)$ to $\mu \in (-\infty, \infty)$.

Given χ , the consumption function can be recovered from

$$\hat{\mathbf{c}} = \underline{\mathbf{c}} + \overbrace{\frac{1}{1 + \exp(-\chi)}}^{=\omega} \Delta h \underline{\kappa}.\tag{13}$$

Thus, the method of moderation is to calculate χ at the points μ corresponding to the log of the Δm points defined above, and then using these to construct an interpolating approximation $\hat{\chi}$ from which we indirectly obtain our approximated consumption rule $\hat{\mathbf{c}}$ (an approximation to the true $\hat{\mathbf{c}}$) by substituting $\hat{\chi}$ for χ in equation (13).

Because this method relies upon the fact that the problem is easy to solve if the decision maker has unreasonable views (either in the optimistic or the pessimistic direction), and because the correct solution is always between these immoderate extremes, we call our solution procedure the ‘method of moderation.’

Results are shown in Figure 3; a reader with very good eyesight might be able to detect the barest hint of a discrepancy between the Truth and the Approximation at the far right-hand edge of the figure, a stark contrast with the calamitous divergence evident in Figure 1.

4. Extensions

4.1. A Tighter Upper Bound

The method described above does not guarantee that the approximated consumption function respects the constraint $\hat{\mathbf{c}}(m) < \bar{\kappa} \Delta m$, where $\bar{\kappa}$ is the MPC at the natural borrowing constraint. Near the constraint, the optimist’s bound $\bar{\mathbf{c}}$ becomes loose because it is calibrated to the low MPC that prevails at high wealth. A tighter upper bound for low-wealth consumers eliminates this slack.

Carroll and Shanker (2024) derives an explicit formula for this maximal MPC: $\bar{\kappa} = 1 - \wp^{1/\rho}(\Phi/R)$ where $\wp$ is the unemployment probability derived by Carroll and Toche (2009), extending the explicit limiting MPC formulas

⁴Under the GIC, $\chi(\mu)$ is asymptotically linear with slope $\alpha \geq 0$ as $\mu \rightarrow +\infty$. We extrapolate χ linearly using the boundary slope, preserving $\omega \in (0, 1)$ and hence $\underline{\mathbf{c}} < \hat{\mathbf{c}} < \bar{\mathbf{c}}$ throughout the extrapolation domain.




Figure 3: Extrapolated $\hat{c}$ Constructed Using the Method of Moderation

established in buffer-stock theory by [Ma and Toda \(2021\)](#). Strict concavity of the consumption function implies $\hat{c}(m) < \bar{\kappa}\Delta m$ for low wealth, while the optimist's bound $\hat{c}(m) < \bar{c}(m) = (\Delta m + \Delta h)\underline{\kappa}$ is tighter for high wealth.




Figure 4: A Tighter Upper Bound

As shown in Figure 4, the two bounds intersect at a cusp point m^* , derived in the Appendix ((?)). Below the cusp we replace the optimist's bound with the tighter one and define a low-resource moderation ratio ((?)), interpolating its logit transform exactly as before. For computational robustness we use a three-piece approximation: the tight bound below the cusp, Hermite interpolation near it (detailed in the Appendix), and the optimist bound above. This yields a continuous, differentiable consumption function respecting all theoretical constraints.

The method of moderation also contributes to literature on improving the

Table 1: Maximum absolute approximation errors by interval. Orders of magnitude in parentheses.

Method	$[m_0, m_1]$	$[m_1, m_2]$	$[m_2, m_3]$	$[m_3, m_4]$	$[m_4, \bar{m}]$
EGM	8.6(-3)	1.8(-4)	2.5(-5)	7.3(-6)	1.1(-1)
MoM	2.9(-3)	4.3(-6)	6.6(-7)	1.3(-7)	2.4(-3)

precision of dynamic stochastic optimization solutions, such as [Chipeniuk \(2020\)](#). Table 1 demonstrates the accuracy gains obtained with the method between each pair of grid points m_j, m_{j+1} , as well as for the extrapolation of the consumption function to $\bar{m} = 30$. Displayed is the maximum absolute difference between the true consumption function and each approximation, evaluated on a dense subgrid of each interval. In each region the method of moderation produces an approximation which is more than an order of magnitude more accurate than the basic EGM.

4.2. Value Function

Often it is useful to know the value function as well as the consumption rule, and the same moderation construction applies. Using the inverse value transformation $\bar{\Lambda} = ((1 - \rho)\bar{\mathbf{v}})^{1/(1-\rho)}$, which is linear in market resources under perfect foresight, we define a value moderation ratio $\bar{\Omega}$ analogous to the consumption moderation ratio, apply the logit transformation, interpolate at gridpoints, and invert to recover $\hat{\mathbf{v}} = \mathbf{u}(\bar{\Lambda})$. The Appendix gives the precise definitions ((??)) and derivation.

4.3. Stochastic Rate of Return

For i.i.d. returns with $\log \mathbf{R} \sim \mathcal{N}(r + \pi - \sigma_{\mathbf{R}}^2/2, \sigma_{\mathbf{R}}^2)$,⁵ [Samuelson \(1969\)](#); [Merton \(1969, 1971\)](#) showed that for a consumer without labor income (or with perfectly forecastable labor income) the consumption function is linear, with an $\text{MPC} = 1 - (\beta \mathbb{E}[\mathbf{R}^{1-\rho}])^{1/\rho}$. See [Carroll \(2020a\)](#); [Benhabib and Bisin \(2018\)](#); [Chipeniuk et al. \(2021\)](#) for extensions. The pessimist and the optimist face certain income but the same stochastic return, so the Merton-Samuelson result applies to both and their consumption functions remain linear. The realist faces both labor income and return risk, and the moderation ratio captures their combined precautionary response. Substitute this stochastic MPC for κ throughout our formulas; the optimist’s and pessimist’s human-wealth definitions $\bar{h}$ and $\underline{h}$ are adjusted in parallel by replacing $\mathbf{R}$ with $\mathbb{E}[\mathbf{R}]$ in their PDV formulas. As Figure 5 shows, consumption remains bounded between the pessimist and the optimist, each of which consume slightly less in the face of return uncertainty.

⁵Here r is the log risk-free rate and π is the equity premium (the expected excess log return). This parametrization ensures $\mathbb{E}[\mathbf{R}] = \exp(r + \pi)$, so that increasing $\sigma_{\mathbf{R}}^2$ constitutes a mean-preserving spread of the level of the return.




Figure 5: Effect of Return Uncertainty on Consumption Bounds

5. Conclusion

The method proposed here is not universally applicable. For example, the method cannot be used for problems for which upper and lower bounds to the ‘true’ solution are not known. But many problems do have obvious upper and lower bounds, and in those cases (as in the consumption example used in the paper), the method may result in substantial improvements in accuracy and stability of solutions. The method of moderation is efficient because the transformed moderation ratio is better-behaved than consumption, requiring fewer gridpoints. As the accuracy results in Table 1 confirm, these gains compound: the method of moderation produces errors more than an order of magnitude smaller than the basic EGM across all grid intervals, including the extrapolation region where standard methods fail most severely.

References

- Aiyagari, S.R., 1994. Uninsured Idiosyncratic Risk and Aggregate Saving. *Quarterly Journal of Economics* 109, 659–684. doi:[10.2307/2118417](https://doi.org/10.2307/2118417).
- Benhabib, J., Bisin, A., 2018. Skewed Wealth Distributions: Theory and Empirics. *Journal of Economic Literature* 56, 1261–1291. doi:[10.1257/jel.20161390](https://doi.org/10.1257/jel.20161390).
- Carroll, C.D., 1997. Buffer-Stock Saving and the Life Cycle/Permanent Income Hypothesis. *Quarterly Journal of Economics* 112, 1–55. doi:[10.1162/003355397555109](https://doi.org/10.1162/003355397555109).
- Carroll, C.D., 2006. The method of endogenous gridpoints for solving dynamic stochastic optimization problems. *Economics Letters* 91, 312–320. doi:[10.1016/j.econlet.2005.09.013](https://doi.org/10.1016/j.econlet.2005.09.013).

- Carroll, C.D., 2009. Precautionary Saving and the Marginal Propensity to Consume Out of Permanent Income. *Journal of Monetary Economics* 56, 780–790. doi:[10.1016/j.jmoneco.2009.06.016](https://doi.org/10.1016/j.jmoneco.2009.06.016).
- Carroll, C.D., 2020a. CRRA utility with rate risk. techreport. Johns Hopkins University. Lecture notes.
- Carroll, C.D., 2020b. Solving microeconomic dynamic stochastic optimization problems. techreport. Johns Hopkins University. URL: <https://www.econ2.jhu.edu/people/ccarroll/SolvingMicroDSOPs.pdf>.
- Carroll, C.D., Kimball, M.S., 1996. On the Concavity of the Consumption Function. *Econometrica* 64, 981–992. doi:[10.2307/2171853](https://doi.org/10.2307/2171853).
- Carroll, C.D., Shanker, A., 2024. Theoretical Foundations of Buffer Stock Saving (Revise and Resubmit, Quantitative Economics). URL: <https://llorracc.github.io/BufferStockTheory/BufferStockTheory.pdf>. revise and Resubmit, Quantitative Economics.
- Carroll, C.D., Toche, P., 2009. A Tractable Model of Buffer Stock Saving. Working Paper 15265. National Bureau of Economic Research. doi:[10.3386/w15265](https://doi.org/10.3386/w15265).
- Chipeniuk, K.O., 2020. Optimal Grid Selection for the Numerical Solution of Dynamic Stochastic Optimization Problems. *Computational Economics* 56, 883–928. doi:[10.1007/s10614-019-09953-4](https://doi.org/10.1007/s10614-019-09953-4).
- Chipeniuk, K.O., Katz, N.H., Walker, T.B., 2021. Households, Auctioneers, and Aggregation. *European Economic Review* 141, 103997. doi:[10.1016/j.euroecorev.2021.103997](https://doi.org/10.1016/j.euroecorev.2021.103997).
- Deaton, A., 1991. Saving and Liquidity Constraints. *Econometrica* 59, 1221–1248. doi:[10.2307/2938366](https://doi.org/10.2307/2938366).
- Friedman, M., 1957. *Theory of the Consumption Function*. Princeton University Press, Princeton, NJ. doi:[10.1515/9780691188485](https://doi.org/10.1515/9780691188485).
- Huggett, M., 1993. The Risk-Free Rate in Heterogeneous-Agent Incomplete-Insurance Economies. *Journal of Economic Dynamics and Control* 17, 953–969. doi:[10.1016/0165-1889\(93\)90024-m](https://doi.org/10.1016/0165-1889(93)90024-m).
- Kimball, M.S., 1990. Precautionary Saving in the Small and in the Large. *Econometrica* 58, 53–73. doi:[10.2307/2938334](https://doi.org/10.2307/2938334).
- Leland, H.E., 1968. Saving and Uncertainty: The Precautionary Demand for Saving. *Quarterly Journal of Economics* 82, 465–473. doi:[10.2307/1879518](https://doi.org/10.2307/1879518).
- Ma, Q., Stachurski, J., Toda, A.A., 2020. The Income Fluctuation Problem and the Evolution of Wealth. *Journal of Economic Theory* 187, 105003. doi:[10.1016/j.jet.2020.105003](https://doi.org/10.1016/j.jet.2020.105003).

- Ma, Q., Toda, A.A., 2021. A Theory of the Saving Rate of the Rich. *Journal of Economic Theory* 192, 105193. doi:[10.1016/j.jet.2021.105193](https://doi.org/10.1016/j.jet.2021.105193).
- Merton, R.C., 1969. Lifetime Portfolio Selection under Uncertainty: The Continuous-Time Case. *Review of Economics and Statistics* 51, 247–257. doi:[10.2307/1926560](https://doi.org/10.2307/1926560).
- Merton, R.C., 1971. Optimum Consumption and Portfolio Rules in a Continuous-Time Model. *Journal of Economic Theory* 3, 373–413. doi:[10.1016/0022-0531\(71\)90038-X](https://doi.org/10.1016/0022-0531(71)90038-X).
- Muth, J.F., 1960. Optimal Properties of Exponentially Weighted Forecasts. *Journal of the American Statistical Association* 55, 299–306. doi:[10.1080/01621459.1960.10482064](https://doi.org/10.1080/01621459.1960.10482064).
- Samuelson, P.A., 1969. Lifetime Portfolio Selection by Dynamic Stochastic Programming. *Review of Economics and Statistics* 51, 239–246. doi:[10.2307/1926559](https://doi.org/10.2307/1926559).
- Sandmo, A., 1970. The Effect of Uncertainty on Saving Decisions. *Review of Economic Studies* 37, 353–360. doi:[10.2307/2296725](https://doi.org/10.2307/2296725).
- Stachurski, J., Toda, A.A., 2019. An Impossibility Theorem for Wealth in Heterogeneous-Agent Models with Limited Heterogeneity. *Journal of Economic Theory* 182, 1–24. doi:[10.1016/j.jet.2019.04.001](https://doi.org/10.1016/j.jet.2019.04.001). provides an explicit linear lower bound on consumption; Corrigendum 2020.
- Zeldes, S.P., 1989. Consumption and Liquidity Constraints: An Empirical Investigation. *Journal of Political Economy* 97, 305–346. doi:[10.1086/261605](https://doi.org/10.1086/261605).